Multilevel/hierarchical models: Overview

Silje Synnøve Lyder Hermansen

2024-02-28

Where are we in the course?

Recap from Monday

When observations are not i.i.d. (i.e. they share a group identity), we will often consider alternatives to the ordinary linear model

▶ negative take: the assumptions of the linear model are not met.
  ▶ non-normal residuals,
  ▶ heteroscedastic residuals
  ▶ correlation between x and residuals
▶ positive take: we have variation that we want to leverage strategically
  ▶ within-group variation
  ▶ between-group variation
  ▶ more correct estimation of the standard errors

⇒ *see this as an opportunity*

I pick my models as part of my research design

What are the most relevant correlations/variation given my theory?

▶ in experiments: you can create that variation and randomize the rest (cut out confounders)

▶ in observational studies: you'll have to "hunt" for the variation you want and control away the rest

Confounders

▶ Control variables that – if absent lead to omitted variable bias – satisfy three criteria:
  ▶ z correlates with y
  ▶ z correlates with x
  ▶ z causes x and y (not intermediate/post-treatment)

  $\rightarrow$ *even when 3 is not satisfied, it might be a sign of a common group identity (e.g. nationality)*

▶ Group identities: observations done in the same context share many potential confounders
  ▶ you might kill several birds with one stone

The principle

The principle

We make the assumption that the residuals are drawn from a normal distribution

▶ **pooled models:** a single distribution

$$y_i = a + bx_i + \epsilon_i$$
$$\epsilon_i \sim N(0, \sigma^2)$$

▶ **hierarchical models:** add a hierarchy
   ▶ assume groups are drawn from different distributions
   ▶ their mean is drawn from a single distribution that "rules them all"

$$y_i = a + bx_i + \epsilon_{ji}$$
$$\epsilon_j \sim N(\alpha_j, \sigma_j^2)$$
$$\alpha_j \sim N(0, \sigma_\alpha^2)$$

Untangling the parameters/variation

This allows me to untangle different sources of variation

$$y_i = a + bx_i + \epsilon_{ji}$$
$$\epsilon_j \sim N(\alpha_j, \sigma_j^2)$$
$$\alpha_j \sim N(0, \sigma_\alpha^2)$$

- ▶ α_j: grouped mean of residuals: group intercept
- ▶ σ_α^2: between-group variation
- ▶ σ_j^2: group-level (within) variation

The promises of a hierarchical structure

This allows me to leverage different sources of variation

- ▶ leverage **within-group** variation:
    - ▶ by factoring out/control for between-group variation (σ_j^2)
- ▶ leverage **between-group** variation:
    - ▶ by running a second regression on the group means (α_α^2)
        - ▶ adjusts the standard errors
        - ▶ data augmentation: add variables from other sources that vary by group
        - ▶ predict out of sample even for new groups
- ▶ leverage **both sources** of variation
   - ▶ by borrowing from the more informative variation ("pooling"/"shrinkage")

Labeling the errors: grouped residuals

Labeling the errors: grouped residuals

Our residuals have group identities that we can "label" as such.

The ghosts of our regression
Residuals correlate with my predictor

Group means and group-level variation

Our residuals have group identities that we can "label" as such.

▶ each group of residuals has a distribution with a mean and a spread

```
## # A tibble: 28 x 3
##    Nationality    y_bar_j sigma2_alpha
##    <chr>            <dbl>        <dbl>
## 1 Austria         -0.665         1.65
## 2 Belgium         -1.54          1.15
## 3 Bulgaria         1.41          2.44
## 4 Croatia          0.549         4.10
## 5 Cyprus          -0.253         1.89
## 6 Czech Republic  -0.206         1.84
## 7 Denmark         -1.48          1.30
## 8 Estonia         -1.33          0.950
## 9 Finland         -1.47          0.919
## 10 France         -1.11          1.26
## # i 18 more rows
```


Residuals can be parsed out as separate distributions

$\Rightarrow$ *I can reconstruct their theoretical distribution by calculating the group mean and standard deviation*

Between-group variation

The group means are drawn from a common normal distribution with a mean and a spread

$\Rightarrow$ *I am treating the residuals as if they were a variable, so statistical theory can be applied*

Varying-intercepts regression: within-group variation

Varying-intercepts regression: within-group variation

The random/varying-intercept model:

- ▶ a common slope for all predictors
- ▶ separate intercepts for all group identities
- ▶ a common intercept (grand mean)

From labelled errors to varying intercepts

Instead of hiding the groupings in the residuals, we can report them as a series of intercepts (i.e. report their group means)

$$y_i = a + bx_i + \alpha_j$$
$$\alpha_j \sim N(0, \sigma_\alpha^2)$$

▶ a: the **grand mean** (mean of α means)
▶ α_j: **varying intercepts** (deviations from this grand mean)

$\Rightarrow$ *useful for interpretation in R*

Varying-intercepts

Now, it is clear that I parse out (control for) between-group variation

▶ **within-group variation** the b coefficients report the effect of observation-level variables

▶ **group-level variation** is reported in the varying intercepts, it is the variation that:
  ▶ has not been accounted for by my main effects
  ▶ that can be attributed to group identities

Estimation in R: Varying national intercepts

Estimation in R: Varying national intercepts

Let's regress MEPs' investment in their district (y) on...

▶ x: their party's size in the national parliament (as a proxy for state funding).
▶ ... while controlling away between-national variation

Equation:

Staff size $= a + b \times$ *Party size* $+ \alpha_{Nationality}$

$y_i = a + bx_i + \alpha_{ij}$

Estimation:

```
library(lme4)
mod.ran.int <- lmer(y ~ x + (1|Nationality),
                    df)
```

Reading the R output

Reading the R output

```
summary(mod.ran.int)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + (1 | Nationality)
##    Data: df
##
## REML criterion at convergence: 31355.2
##
## Scaled residuals:
##     Min      1Q  Median      3Q     Max
## -3.1127 -0.5387 -0.1435  0.3598 15.2357
##
## Random effects:
##  Groups      Name        Variance Std.Dev.
##  Nationality (Intercept) 3.125    1.768
##  Residual                5.240    2.289
## Number of obs: 6948, groups: Nationality, 28
##
## Fixed effects:
##             Estimate Std. Error t value
## (Intercept)   2.6799     0.3386   7.915
## x            -1.6722     0.1678  -9.965
##
## Correlation of Fixed Effects:
##   (Intr)
## x -0.117
```

R refers to the residuals as "random effects"

σ_α^2: remaining **between-group variance**: 3.12

- standard deviation: 1.77
- the unexplained variation between groups

Residual: remaining **within-group variance**: 5.24

- standard deviation of within-group distribution: 2.29
- the unexplained variation in all groups

R refers to regression coefficients as "fixed effects"

a: intercept/**grand mean**: 2.68

- a hypothetical intercept for interpretation (mean of means)

b: **slope**: -1.67

- the marginal effect of party size (x)

Interpretation

Interpretation

Interpretation follows normal principles, but there are some complications:

a. we now have two intercepts per scenario:

▶ the grand mean (a): for focus on general effect of x
▶ the group-level mean (α_j): for description and prediction
▶ the grand mean (a): for focus on general effect of x
▶ sum of the grand mean (a) group-level mean (α_j): for prediction

b. all effects are linear

▶ so first-difference and marginal effects are the same

Interpreting marginal effects

The interpretation of the marginal effect is as with any linear model:

Table 1: Effect of state funding for parties on MEPs' local staff size

	Dependent variable:
	y
x	-1.672^{***}
	(0.168)
Constant	2.680^{***}
	(0.339)
Observations	6,948
Log Likelihood	$-15,677.610$
Akaike Inf. Crit.	31,363.210
Bayesian Inf. Crit.	31,390.600
Note:	*p<0.1; **p<0.05; ***p<0.01

$\Rightarrow$ *A 10% decrease in the national party's seat share would lead every 6th MEP to compensate by hiring an additional local staffer.*

Prediction

The varying intercepts are reported as deviations from the grand mean

```
fixef(mod.ran.int); ranef(mod.ran.int)
```

```
## (Intercept)        x
##    2.679887  -1.672226
```

```
##                 (Intercept)
## Austria         -0.49518857
## Belgium         -1.52249566
## Bulgaria         1.54657524
## Croatia          0.68267309
## Cyprus          -0.05313986
## Czech Republic  -0.10587832
```

Predicted local staff in Austria when national party is not in Parliament:

▶ $2.68 + \text{-}0.5 \times 0 = 2.18$

Predicted local staff in Austria when national party holds 10% of the seats

▶ $2.68 + \text{-}0.5 + \text{-}1.67 \times 0.1 = 2.02$

Visualization

Varying intercepts, fixed slope

$\Rightarrow$ *the slope is constant, but the intercept changes across nationalities*

Varying slopes, varying intercepts

Defintion

We can let the effect of x vary by group through an interaction effect

$$y_i = a + bx_i + c_j z_i + \alpha_{ji}$$

- ▶ c_j: **varying slope** (the effect of z varies by group)
- ▶ α_j: **varying intercepts**

$\Rightarrow$ *a series of regressions within the regression*

Estimation in R

▶ the estimation is done as if it was an interaction effect
▶ the R syntax is somewhat different

```
mod.ran.slope <- lmer(y ~ x + (ProxNatElection | Nationality),
```

Interpretation

Marginal effects

We can read these coefficients as if they were from separate models

```
ranef(mod.ran.slope)
```

```
##               (Intercept) ProxNatElection
## Austria        -0.3201686     0.001073577
## Belgium        -1.3944093    -0.018299157
## Bulgaria        1.8027887     0.086786759
## Croatia         0.9352286     0.058233060
## Cyprus          0.1020429    -0.003477477
## Czech Republic  0.1112017     0.024050889
```

MEPs from Austria hire on average 0.004 assistants more immediately before an election compared to immediately after, while MEPs from Belgium hire on average 0.073 fewer assistants.

► These are negligible marginal effects.

Prediction

The prediction is done per group, but follows normal rules

- ▶ two intercepts:

  grand mean + group-level intercept

- ▶ one slope per group

Austria after election:

- ▶ $2.51 + -0.32 + 0.001 \times -4 = 2.189$

Austria before election:

- ▶ $2.51 + -0.32 + 0.001 \times 0 = 2.193$

```
fixef(mod.ran.slope); ranef(mod.ran.slope)
```

```
## (Intercept)          x
##   2.513035   -1.691321
```

```
##                (Intercept) ProxNatElection
## Austria         -0.3201686     0.001073577
## Belgium         -1.3944093    -0.018299157
## Bulgaria         1.8027887     0.086786759
## Croatia          0.9352286     0.058233060
## Cyprus           0.1020429    -0.003477477
## Czech Republic   0.1112017     0.024050889
```

Visualization

Varying slopes and intercepts: Effect of electoral calendar by nationality

Level-2 regression: between-group variation

Definition

Definition

We can think of the residuals/group intercepts as a variable in their own right

$$y_i = a + bx_i + \epsilon_{ji}$$

▶ they are generated by draws from J number of distributions: $\epsilon_{ji} \sim N(\alpha_j, \sigma_\alpha^2)$

▶ ... and therefore we can model them

$$\alpha_j = c_1 + c_2 z_j$$

$\Rightarrow$ *we run a second regression on the residuals*

Implications

We explicitly model between-group variation

- ▶ z, the level-2 predictor only varies at the group level
    - ▶ standard errors for z reflect the number of groups
    - ▶ the more groups, the more the approach makes sense
- ▶ data augmentation
    - ▶ we can add information from other to the model
    - ▶ contextual elements
    - ▶ improves prediction

Estimation in R: Electoral system

Estimation in R: Electoral system

Let's add electoral system (z) as a predictor

▶ it never changes in a country (in this study)

R handles this automatically

▶ same data frame
  ▶ all variables that don't vary within groups are regressed as a level 2
▶ coefficients reported the same way
▶ estimation of coefficients and standard errors is different

```
mod.two.levels <- lmer(y ~ x + z + (1|Nationality), df)
```

Reading the R output

Reading the R output

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ x + z + (1 | Nationality)
##    Data: df
##
## REML criterion at convergence: 31353.9
##
## Scaled residuals:
##     Min      1Q  Median      3Q     Max
## -3.1145 -0.5388 -0.1434  0.3599 15.2339
##
## Random effects:
##  Groups      Name        Variance Std.Dev.
##  Nationality (Intercept) 3.235    1.799
##  Residual                5.240    2.289
## Number of obs: 6948, groups: Nationality, 28
##
## Fixed effects:
##             Estimate Std. Error t value
## (Intercept)   2.5268     0.6030   4.191
## x            -1.6719     0.1678  -9.962
## z             0.2263     0.7311   0.310
##
## Correlation of Fixed Effects:
##   (Intr) x
## x -0.077
## z -0.821  0.013
```

The level-2 regression coefficient appears as "fixed effects"

c: **slope:** 0.23

► the marginal effect of electoral system (z)

Check the **change in between-group variance**:

► the between-group variance (σ_α^2, 3.23) should decrease
► it is not the case here (3.12 $\leq$ 3.23)

$\rightarrow$ *increase in variance indicates "complexities" between levels (interactions)*

Correlation of Fixed Effects:

► negative correlation between predictor (z) and intercept (-0.82): high level of z correlates with low base-line value of y.

Silje Synnøve Lyder Hermansen   Multilevel/hierarchical models: Overview   2024-02-28   41 / 41