

Missing data

Silje Synnøve Lyder Hermansen

Recap: our course

Recap: our course

We are entering the last part of this course

1. R-skills, interpretation and non-linear effects
2. Data structures
 - ▶ Hierarchical data
3. Limited and categorical outcome variables (GLMs)
4. **Data structures** (today/last week, short)
 - ▶ **Missing data**

The purpose of this course

The purpose of this course

⇒ The purpose of this course is to find solutions when the assumptions of the linear model are not satisfied

Two assumptions in ordinary regression

Two assumptions in ordinary regression

Linear models (OLS) rely on two assumptions that are often violated

1. **Assumption 1:** outcomes are continuous and unbounded
2. **Assumption 2:** observations are independent and identically distributed
 - ▶ independent: probability of observing one unit is not dependent on observing another
 - ▶ identically distributed: observations come from the same data generating process/probability distribution

⇒ *strategies for when these are not satisfied*

Solutions to violations of those assumptions

1. Assumption 1: Limited and categorical outcome variables (GLMs): - recode the dependent variable and describe the data generating process w/probability distribution - choice of model depends on the data generating process - e.g. logit, multinomial, ordinal, poisson, neg.bin, zero-inflated, coxph...

2. Assumption 2: Observations are not iid: - hierarchical/nested data - *missing data*

⇒ *what do we do when observations are not iid?*

Today (week 13 and 14)

Sources of missing data

Sources of missing data

Most data contain missing observations

- ▶ missing data (NA) is the result of a “lurking” variable that :
 - ▶ assigns NA to some of the other variables
 - ▶ ... possibly affecting both x and y
- ▶ the “lurking” means that the assignment mechanism is not observed
 - ▶ think about the data generating process of the NA
 - ▶ we have to theorize/make assume

⇒ *missing observations might indicate a confounder*

What to do?

What to do?

$\Rightarrow$ *addressing/reducing the problem is often easier than what we think*

Classifications of missing data:

Take 1

The original classification by Rubin (1979)

- ▶ **MCAR** (Missing Completely at Random)
 - ▶ probability of NA is the same for all cases
- ▶ **MAR** (Missing at Random):
 - ▶ probability of NA depends on *observable* data (known sources)
- ▶ **MNAR** (Missing Not at Random)
 - ▶ probability of NA depends on *unobservable* data (unknown sources)

⇒ *these are assumptions that we can never test*

Why is it a problem

- ▶ **statistical power (MCAR):** only a problem if it reduces the N too much → *a representative sample*
- ▶ **information bias (MAR, MNAR):** we only record parts of a phenomenon (recall bias, missclassification, observer bias. . .)
 - ▶ independent variables:
 - ▶ we might not get the full “elasticity” of the variable
 - ▶ dependent variable: do we underestimate our phenomenon?
- ▶ **selection bias (MAR, MNAR):**
 - ▶ our estimate is biased because the unobserved assignment of NA affects both x and y

Take 2

We can subdivide the last category

- ▶ **MCAR** (Missing Completely at Random)
 - ▶ NA are *not* dependent on any predictors (observed or not): not conditional → *you can ignore the problem, unless you have too little statistical power*
 - ▶ **MAR** (Missing at Random):
 - ▶ NA depends on the value of other observed predictors: it is conditional → *ignorability; you can condition on the other predictors*
 - ▶ **MNAR** (Missing Not at Random)
 - ▶ NA depends on unobserved data
 - ▶ NA depends on the value of *the predictor itself* (e.g. censoring)
- *NA must be modeled, or you will have to accept a biased estimate*

Strategies

Simple strategies

Discard data

Ignore the problem

- ▶ complete case analysis:
 - ▶ the usual “listwise.exclusion”
- ▶ available data analysis:
 - ▶ analyze subsets of data separately
 - ▶ exclude variables with missing observations
- ▶ weighing of NA according to predictors
 - ▶ common in surveys → *some cases may be underrepresented in the data, because of NA*

Replace each NA by a single value

We can also infer the missing values in fairly simple ways

- ▶ **mean imputation:**
 - ▶ replace the missing data by the variable mean
- ▶ **conditional mean imputation:** use information from other variables
 - ▶ group mean, regression predictions

⇒ *doesn't take into account the uncertainty from our estimate*

Multiple imputations

Multiple imputations

Multiple imputation generates several predictions for each missing value to account for the uncertainty

- ▶ step 1: make predictions for the missing values by adding some random noise for each model

→ *we end up with several data sets (5-20 frames)*

- ▶ step 2: estimate the main model on all the different data sets

→ *pool over the regression parameters*

EM algorithm

EM algorithm

The EM is the base-line approach, and only has one data frame in the end

- ▶ we have several variables
- ▶ E-step:
 - ▶ give your NA some initial values
 - ▶ predict your x_{miss} using the observed values and initial values of x (and all other predictors)
- ▶ M-step:
 - ▶ re-do until your predictions of x_{miss} don't change any more (set a value at which you stop)

⇒ *classic maximum likelihood with a twist*

Multiple Imputation via Chained Equations (MICE)

Multiple Imputation via Chained Equations (MICE)

We assume a set of variables are correlated, and use them to predict for each other in turn (a cycle)

Imagine x , y and z :

- ▶ cycle 1:
 - ▶ $x \propto \alpha + \beta_1 y + \beta_2 z$: give y and z some starting value; regress x on all other models
 - ▶ $y \propto \alpha + \beta_1 \hat{x} + \beta_2 z$: replace missing values of x by predicted $\hat{x}$
 - ▶ $z \propto \alpha + \beta_1 \hat{x} + \beta_2 \hat{y}$: same
- ▶ cycle 2-...: rinse and repeat until nothing changes (convergence)



Figure 1: Mice thrive in holes...

... and add some random noise

This is usually done together with a bit of (random) noise at each step

- ▶ for each iteration, create a new data set with imputed variables
- ▶ run regular (g)lm on each data set:
 - ▶ regression parameters (β) are averaged over
 - ▶ the standard error is a fusion of both:
 - ▶ within-model variation: standard errors from the regression
 - ▶ between-model variation: the deviation between the regression parameters

..then check

... were my imputations appropriate?

- ▶ problem of **overfitting**:
 - ▶ you may have perfect in-sample predictions, that are useless for out-of-sample imputation
- ▶ **overimputation**: randomly leave observations out, check if you predict correctly

Literature